



AAKASH DESHMANE \

133008022 \

ECEN 743 ASSIGNMENT 4 \

ECEN 743: Reinforcement Learning

Deep Q-Learning

Code tested using

1. gymnasium 0.27.1
2. box2d-py 2.3.5
3. pytorch 2.0.0
4. Python 3.9.12

1 & 2 can be installed using `pip install gymnasium[box2d]`

General Instructions

1. This code consists of TODO blocks, read them carefully and complete the blocks

2. Type your code between the following lines

TYPE YOUR CODE HERE

#####

3. The default hyperparameters should be able to solve LunarLander-v2

4. You do not need to modify the rest of the code for this assignment, free to do so if needed.

AAKASH DESHMANE

133008022

ECEN 743 ASSIGNMENT 4

ECEN 743: Reinforcement Learning Deep Q-Learning Code tested using 1. gymnasium 0.27.1 2. box2d-py 2.3.5 3. pytorch 2.0.0 4. Python 3.9.12 1 & 2 can be installed using `pip install gymnasium[box2d]` General Instructions

1. This code consists of TODO blocks, read them carefully and complete each of the blocks

2. Type your code between the following lines

TYPE YOUR CODE HERE

#####

3. The default hyperparameters should be able to solve LunarLander-v2

4. You do not need to modify the rest of the code for this assignment, feel free to do so if needed.



```
1 !pip install gymnasium[all]
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: gymnasium[all] in /usr/local/lib/python3.9/dist-packages (0.28.1)
Requirement already satisfied: farama-notifications>=0.0.1 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (0.0.4)
Requirement already satisfied: importlib-metadata>=4.8.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (6.1.0)
Requirement already satisfied: typing-extensions>=4.3.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (4.5.0)
Requirement already satisfied: cloudpickle>=1.2.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (2.2.1)
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (1.22.4)
Requirement already satisfied: jax-jumpy>=1.0.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (1.0.0)
Requirement already satisfied: box2d-py==2.3.5 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (2.3.5)
Requirement already satisfied: imageio>=2.14.1 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (2.25.1)
Requirement already satisfied: mujoco-py<2.2,>=2.1 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (2.1.2.14)
Requirement already satisfied: opencv-python>=3.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (4.7.0.72)
Requirement already satisfied: pygame==2.1.3 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (2.1.3)
Requirement already satisfied: jax==0.3.24 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (0.3.24)
Requirement already satisfied: moviepy>=1.0.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (1.0.3)
Requirement already satisfied: shimmy[atari]<1.0,>=0.1.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (0.2.1)
Requirement already satisfied: jaxlib==0.3.24 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (0.3.24)
Requirement already satisfied: torch>=1.0.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (2.0.0+cu118)
Requirement already satisfied: mujoco>=2.3.2 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (2.0.3.3)
Requirement already satisfied: swig==4.* in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (4.1.1)
Requirement already satisfied: matplotlib>=3.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (3.7.1)
Requirement already satisfied: lz4>=3.1.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[all]) (4.3.2)
Requirement already satisfied: opt-einsum in /usr/local/lib/python3.9/dist-packages (from jax==0.3.24->gymnasium[all]) (3.3.0)
Requirement already satisfied: scipy>=1.5 in /usr/local/lib/python3.9/dist-packages (from jax==0.3.24->gymnasium[all]) (1.10.1)
Requirement already satisfied: pillow>=8.3.2 in /usr/local/lib/python3.9/dist-packages (from imageio>=2.14.1->gymnasium[all]) (8.4.0)
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.9/dist-packages (from importlib-metadata>=4.8.0->gymnasium[all]) (0.8.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.9/dist-packages (from matplotlib>=3.0->gymnasium[all]) (3.0.9)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.9/dist-packages (from matplotlib>=3.0->gymnasium[all]) (1.0.7)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.9/dist-packages (from matplotlib>=3.0->gymnasium[all]) (1.4.5)
Requirement already satisfied: importlib-resources>=3.2.0 in /usr/local/lib/python3.9/dist-packages (from matplotlib>=3.0->gymnasium[all]) (5.12.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.9/dist-packages (from matplotlib>=3.0->gymnasium[all]) (0.11.0)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.9/dist-packages (from matplotlib>=3.0->gymnasium[all]) (2.8.2)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.9/dist-packages (from matplotlib>=3.0->gymnasium[all]) (4.22.0)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.9/dist-packages (from matplotlib>=3.0->gymnasium[all]) (23.1)
Requirement already satisfied: requests<3.0,>=2.8.1 in /usr/local/lib/python3.9/dist-packages (from moviepy>=1.0.0->gymnasium[all]) (2.28.1)
Requirement already satisfied: proglog<=1.0.0 in /usr/local/lib/python3.9/dist-packages (from moviepy>=1.0.0->gymnasium[all]) (0.1.10)
Requirement already satisfied: imageio-ffmpeg>=0.2.0 in /usr/local/lib/python3.9/dist-packages (from moviepy>=1.0.0->gymnasium[all]) (0.2.0)
Requirement already satisfied: tqdm<5.0,>=4.11.2 in /usr/local/lib/python3.9/dist-packages (from moviepy>=1.0.0->gymnasium[all]) (4.65.0)
Requirement already satisfied: decorator<5.0,>=4.0.2 in /usr/local/lib/python3.9/dist-packages (from moviepy>=1.0.0->gymnasium[all]) (4.4.2)
Requirement already satisfied: glfw in /usr/local/lib/python3.9/dist-packages (from mujoco>=2.3.2->gymnasium[all]) (2.5.9)
Requirement already satisfied: absl-py in /usr/local/lib/python3.9/dist-packages (from mujoco>=2.3.2->gymnasium[all]) (1.4.0)
Requirement already satisfied: pyopengl in /usr/local/lib/python3.9/dist-packages (from mujoco>=2.3.2->gymnasium[all]) (3.1.6)
Requirement already satisfied: cffi>=1.10 in /usr/local/lib/python3.9/dist-packages (from mujoco-py<2.2,>=2.1->gymnasium[all]) (1.15.1)
Requirement already satisfied: Cython>=0.27.2 in /usr/local/lib/python3.9/dist-packages (from mujoco-py<2.2,>=2.1->gymnasium[all]) (0.29.36)
Requirement already satisfied: fasteners>=0.15 in /usr/local/lib/python3.9/dist-packages (from mujoco-py<2.2,>=2.1->gymnasium[all]) (0.17.3)
Requirement already satisfied: ale-py>=0.8.1 in /usr/local/lib/python3.9/dist-packages (from shimmy[atari]<1.0,>=0.1.0->gymnasium[all]) (0.8.1)
Requirement already satisfied: sympy in /usr/local/lib/python3.9/dist-packages (from torch>=1.0.0->gymnasium[all]) (1.11.1)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.9/dist-packages (from torch>=1.0.0->gymnasium[all]) (3.1.2)
Requirement already satisfied: triton==2.0.0 in /usr/local/lib/python3.9/dist-packages (from torch>=1.0.0->gymnasium[all]) (2.0.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.9/dist-packages (from torch>=1.0.0->gymnasium[all]) (3.10.7)
Requirement already satisfied: networkx in /usr/local/lib/python3.9/dist-packages (from torch>=1.0.0->gymnasium[all]) (3.0)
```

```
Requirement already satisfied: cmake in /usr/local/lib/python3.9/dist-packages (from triton==2.0.0->torch>=1.0.0->gymnasium[all]) (3.
Requirement already satisfied: lit in /usr/local/lib/python3.9/dist-packages (from triton==2.0.0->torch>=1.0.0->gymnasium[all]) (16.0
Requirement already satisfied: pycparser in /usr/local/lib/python3.9/dist-packages (from cffi>=1.10-> mujoco-py<2.2,>=2.1->gymnasium[a
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.9/dist-packages (from python-dateutil>=2.7->matplotlib>=3.0->gymnas
```

```
1 !pip3 install torch
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: torch in /usr/local/lib/python3.9/dist-packages (2.0.0+cu118)
Requirement already satisfied: filelock in /usr/local/lib/python3.9/dist-packages (from torch) (3.10.7)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.9/dist-packages (from torch) (4.5.0)
Requirement already satisfied: networkx in /usr/local/lib/python3.9/dist-packages (from torch) (3.0)
Requirement already satisfied: triton==2.0.0 in /usr/local/lib/python3.9/dist-packages (from torch) (2.0.0)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.9/dist-packages (from torch) (3.1.2)
Requirement already satisfied: sympy in /usr/local/lib/python3.9/dist-packages (from torch) (1.11.1)
Requirement already satisfied: cmake in /usr/local/lib/python3.9/dist-packages (from triton==2.0.0->torch) (3.25.2)
Requirement already satisfied: lit in /usr/local/lib/python3.9/dist-packages (from triton==2.0.0->torch) (16.0.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.9/dist-packages (from Jinja2->torch) (2.1.2)
Requirement already satisfied: mpmath>=0.19 in /usr/local/lib/python3.9/dist-packages (from sympy->torch) (1.3.0)
```

```
1 import gymnasium as gym
2 import random
3 import torch
4 import torch.nn as nn
5 import torch.nn.functional as F
6 import torch.optim as optim
7 import argparse
8 import numpy as np
9 from collections import deque, namedtuple
10 import matplotlib.pyplot as plt
11 import base64, io
12 import cv2
13
14
```

```
/usr/local/lib/python3.9/dist-packages/ipykernel/ipkernel.py:283: DeprecationWarning: `should_run_async` will not call `transform_cell`
and should_run_async(code)
```

```
1 class ExperienceReplay:
2     """
3     Based on the Replay Buffer implementation of TD3
4     Reference: https://github.com/sfujim/TD3/blob/master/utils.py
5     """
6     def __init__(self, state_dim, action_dim, max_size, batch_size, gpu_index=0):
7         self.max_size = max_size
8         self.ptr = 0
9         self.size = 0
10        self.state = np.zeros((max_size, state_dim))
11        self.action = np.zeros((max_size, action_dim))
12        self.next_state = np.zeros((max_size, state_dim))
13        self.reward = np.zeros((max_size, 1))
14        self.done = np.zeros((max_size, 1))
15        self.batch_size = batch_size
16        self.device = torch.device('cuda', index=gpu_index) if torch.cuda.is_available() else torch.device('cpu')
17
18
19    def add(self, state, action, reward, next_state, done):
20        self.state[self.ptr] = state
21        self.action[self.ptr] = action
22        self.next_state[self.ptr] = next_state
23        self.reward[self.ptr] = reward
24        self.done[self.ptr] = done
25        self.ptr = (self.ptr + 1) % self.max_size
26        self.size = min(self.size + 1, self.max_size)
27
28    def sample(self):
29        ind = np.random.randint(0, self.size, size=self.batch_size)
30
31        return (
32            torch.FloatTensor(self.state[ind]).to(self.device),
33            torch.FloatTensor(self.action[ind]).long().to(self.device),
34            torch.FloatTensor(self.reward[ind]).to(self.device),
35            torch.FloatTensor(self.next_state[ind]).to(self.device),
36            torch.FloatTensor(self.done[ind]).to(self.device)
37        )
```

38
39

```

1 class QNetwork(nn.Module):
2     """
3     Q Network: designed to take state as input and give out Q values of actions as output
4     """
5
6     def __init__(self, state_dim, action_dim):
7         """
8         state_dim (int): state dimension
9         action_dim (int): action dimension
10        """
11        super(QNetwork, self).__init__()
12        self.l1 = nn.Linear(state_dim, 64)
13        self.l2 = nn.Linear(64, 64)
14        self.l3 = nn.Linear(64, action_dim)
15
16    def forward(self, state):
17        q = F.relu(self.l1(state))
18        q = F.relu(self.l2(q))
19        return self.l3(q)
20
21
22 class DQNAgent():
23
24     def __init__(self,
25         state_dim,
26         action_dim,
27         discount=0.99,
28         tau=1e-3,
29         lr=5e-4,
30         update_freq=4,
31         max_size=int(1e5),
32         batch_size=64,
33         gpu_index=0
34     ):
35         """
36         state_size (int): dimension of each state
37         action_size (int): dimension of each action
38         discount (float): discount factor
39         tau (float): used to update q-target
40         lr (float): learning rate
41         update_freq (int): update frequency of target network
42         max_size (int): experience replay buffer size
43         batch_size (int): training batch size
44         gpu_index (int): GPU used for training
45        """
46        self.state_dim = state_dim
47        self.action_dim = action_dim
48        self.discount = discount
49        self.tau = tau
50        self.lr = lr
51        self.update_freq = update_freq
52        self.max_size = max_size
53        self.batch_size = batch_size
54        self.device = torch.device('cuda', index=gpu_index) if torch.cuda.is_available() else torch.device('cpu')
55
56        # Setting up the NNs
57        self.Q = QNetwork(state_dim, action_dim).to(self.device)
58        self.Q_target = QNetwork(state_dim, action_dim).to(self.device)
59        self.optimizer = optim.Adam(self.Q.parameters(), lr=self.lr)
60
61        # Experience Replay Buffer
62        self.memory = ExperienceReplay(state_dim,1,max_size,self.batch_size,gpu_index)
63
64        self.t_train = 0
65
66    def step(self, state, action, reward, next_state, done):
67        """
68        1. Adds (s,a,r,s') to the experience replay buffer, and updates the networks
69        2. Learns when the experience replay buffer has enough samples
70        3. Updates target network
71        """
72        self.memory.add(state, action, reward, next_state, done)
73        self.t_train += 1

```

```

54
55     # Experience Replay
56     if self.memory.size > self.batch_size:
57         experiences = self.memory.sample()
58         self.learn(experiences, self.discount) #To be implemented
59
60     # Target Network
61     if (self.t_train % self.update_freq) == 0:
62         self.target_update(self.Q, self.Q_target, self.tau) #To be implemented
63
64 def select_action(self, state, epsilon = 0.):
65
66     if np.random.random() > epsilon:
67         state = torch.tensor(state, dtype=torch.float32, device=self.device) # converting our state to pytorch tensor
68         self.Q.eval()
69         with torch.no_grad():
70             actions=self.Q.forward(state)
71         self.Q.train()
72         action = torch.argmax(actions).item() # .item() converts tensors to integers
73     else:
74         action=np.random.choice(self.action_dim)
75
76     return action
77
78
79 def learn(self, experiences, discount):
80     """
81     TODO: Complete this block to update the Q-Network using the target network
82     1. Compute target using self.Q_target ( target = r + discount * max_b [Q_target(s,b)] )
83     2. Compute Q(s,a) using self.Q
84     3. Compute MSE loss between step 1 and step 2
85     4. Update your network
86     Input: experiences consisting of states,actions,rewards,next_states and discount factor
87     Return: None
88     """
89     states, actions, rewards, next_states, dones = experiences
90
91     q_eval = self.Q(states).gather(1, actions)
92     q_next = self.Q_target(next_states).detach().max(1)[0].unsqueeze(1)
93
94     # Calculating q_target
95     q_target = rewards + discount * q_next * (1-dones)
96
97     # Calculating loss and backpropogating
98     loss = F.mse_loss(q_eval, q_target)
99     self.optimizer.zero_grad()
100    loss.backward()
101    self.optimizer.step()
102    self.target_update(self.Q, self.Q_target, tau)
103
104 def target_update(self, Q, Q_target, tau):
105     """
106     TODO: Update the target network parameters (param_target) using current Q parameters (param_Q)
107     Perform the update using tau, this ensures that we do not change the target network drastically
108     1. param_target = tau * param_Q + (1 - tau) * param_target
109     Input: Q,Q_target,tau
110     Return: None
111     """
112     ##### TYPE YOUR CODE HERE #####
113     for param_target, param_local in zip(Q_target.parameters(), Q.parameters()):
114         param_target.data.copy_(tau*param_local.data + (1.0-tau)*param_target.data)
115
116     #param_target = tau * param_Q + (1 - tau) * param_target
117
11
12 if __name__ == "__main__":
13     seed = 0
14     n_episodes = 1500
15     batch_size = 64
16     discount = 0.99
17     lr = 5e-4 # learning rate
18     tau = 0.001 # soft update of target network
19     max_size = int(1e5)
20     update_freq = 4
21     gpu_index = 0

```

```

12 max_eps_len = 1000
13 #exploration strategy
14 epsilon_start = 1          # start value of epsilon
15 epsilon_end = 0.01        # end value of epsilon
16 epsilon_decay = 0.995     # decay value of epsilon
17
18
19 # making the environment
20 env = gym.make("LunarLander-v2", render_mode="rgb_array")
21
22 #setting seeds
23 torch.manual_seed(seed)
24 np.random.seed(seed)
25 random.seed(seed)
26
27 state_dim = env.observation_space.shape[0]
28 action_dim = env.action_space.n
29
30 kwargs = {
31     "state_dim":state_dim,
32     "action_dim":action_dim,
33     "discount":discount,
34     "tau":tau,
35     "lr":lr,
36     "update_freq":update_freq,
37     "max_size":max_size,
38     "batch_size":batch_size,
39     "gpu_index":gpu_index
40 }
41 learner = DQNAgent(**kwargs) #Creating the DQN learning agent
42 moving_window = deque(maxlen=100)
43 scores=[]
44 state_dict=[]
45 best_state=0
46 index=0
47 epsilon_by_step = lambda step: float(epsilon_end+(epsilon_start - epsilon_end)*np.exp(-1. * step / epsilon_decay))
48
49 # For visualization
50 fourcc = cv2.VideoWriter_fourcc(*'mp4v')
51 video = cv2.VideoWriter('lunar_landerv1.mp4', fourcc, 30, (600, 400))
52 RENDER= True
53 max_array= -1000000
54
55 # Training code
56 for e in range(n_episodes):
57     state, _ = env.reset(seed=seed)
58     curr_reward = 0
59     for t in range(max_eps_len):
60         epsilon = epsilon_by_step(t)
61         action = learner.select_action(state, epsilon) #To be implemented
62         n_state,reward,terminated,truncated,_ = env.step(action)
63         done = terminated or truncated
64         learner.step(state, action, reward, n_state, done) #To be implemented
65         state = n_state
66         curr_reward += reward
67         #if RENDER:
68             #frame = env.render()
69
70             #video.write(frame)
71         if done:
72             break
73     moving_window.append(curr_reward)
74     scores.append(curr_reward)
75
76     if np.mean(moving_window) > 200.00:
77         torch.save(agent.Q.state_dict(), 'checkpoint.pth')
78
79     #if RENDER:
80         #frame = env.render()
81         #frame = cv2.resize(frame, (600,400))
82         #video.write(frame)
83
84
85     """
86     TODO: Write code for decaying the exploration rate using args.epsilon_decay
87     and args.epsilon_end. Note that epsilon has been initialized to args.epsilon_start
88     1. You are encouraged to try new methods

```

```


89     """
90     ##### TYPE YOUR CODE HERE #####
91     #####
92
93     if e % 100 == 0:
94         print('Episode Number {} Average Episodic Reward (over 100 episodes): {:.2f}'.format(e, np.mean(moving_window)))
95     #if curr_reward >= scores.max():
96         #best_state=state_dict(e)
97         #index=e
98
99     """
100    TODO: Write code for
101    1. Logging and plotting
102    2. Rendering the trained agent
103    """
104    fig = plt.figure()
105    ax = fig.add_subplot(111)
106    plt.plot(np.arange(len(scores)), scores)
107    plt.ylabel('Score')
108    plt.xlabel('Episode #')
109    plt.show()
110
111    #video.release()
112    ##### TYPE YOUR CODE HERE #####

```

```

Episode Number 0 Average Episodic Reward (over 100 episodes): -121.30
Episode Number 100 Average Episodic Reward (over 100 episodes): -151.32
Episode Number 200 Average Episodic Reward (over 100 episodes): -34.27
Episode Number 300 Average Episodic Reward (over 100 episodes): 56.14
Episode Number 400 Average Episodic Reward (over 100 episodes): 222.30
Episode Number 500 Average Episodic Reward (over 100 episodes): 196.65
Episode Number 600 Average Episodic Reward (over 100 episodes): 261.28
Episode Number 700 Average Episodic Reward (over 100 episodes): 285.49
Episode Number 800 Average Episodic Reward (over 100 episodes): 260.90
Episode Number 900 Average Episodic Reward (over 100 episodes): 240.62
Episode Number 1000 Average Episodic Reward (over 100 episodes): 265.98
Episode Number 1100 Average Episodic Reward (over 100 episodes): 267.61
Episode Number 1200 Average Episodic Reward (over 100 episodes): 257.59
Episode Number 1300 Average Episodic Reward (over 100 episodes): 254.96
Episode Number 1400 Average Episodic Reward (over 100 episodes): 237.16

```



```

1 # For visualization
2 from gym.wrappers.monitoring import video_recorder
3 from IPython.display import HTML
4 from IPython import display
5 import glob

1 def show_video(env_name):
2     mp4list = glob.glob('video/*.mp4')
3     if len(mp4list) > 0:
4         mp4 = 'video/{}.mp4'.format(env_name)
5         video = io.open(mp4, 'r+b').read()
6         encoded = base64.b64encode(video)
7         display.display(HTML(data='''<video alt="test" autoplay

```

```

8         loop controls style="height: 400px;">
9         <source src="data:video/mp4;base64,{0}" type="video/mp4" />
10        </video>'''.format(encoded.decode('ascii'))))
11    else:
12        print("Could not find video")
13
14 def show_video_of_model(agent, env_name):
15     env = gym.make(env_name, render_mode="rgb_array")
16     fourcc = cv2.VideoWriter_fourcc(*'mp4v')
17     video = cv2.VideoWriter('lunar_lander.mp4', fourcc, 30, (600, 400))
18     agent.Q.load_state_dict(torch.load('checkpoint.pth'))
19     state, _ = env.reset()
20     print(state.shape)
21     done = False
22     while not done:
23         frame = env.render()
24         video.write(frame)
25
26         action = agent.select_action(state)
27
28         n_state, reward, terminated, truncated, _ = env.step(action)
29         done = terminated or truncated
30         agent.step(state, action, reward, n_state, done) #To be implemented
31         state = n_state
32     env.close()
33     video.release()

```



```

1 state, _ = env.reset()

```



```

1 agent = DQNAgent(state_dim=8, action_dim=4)
2 show_video_of_model(agent, 'LunarLander-v2')
3

```



```

(8,)

```