

ECEN 743: Reinforcement Learning Deep Q-Learning Code tested using 1. gymnasium 0.27.1 2. box2d-py 2.3.5 3. pytorch 2.0.0 4. Python 3.9.12

1 & 2 can be installed using `pip install gymnasium[box2d]` General Instructions

1. This code consists of TODO blocks, read them carefully and complete each of the blocks
2. Type your code between the following lines

```
##### TYPE YOUR CODE HERE #####
#####
```

3. The default hyperparameters should be able to solve LunarLander-v2
4. You do not need to modify the rest of the code for this assignment, feel free to do so if needed.

```
1 !pip install gymnasium[box2d]
2 !pip3 install matplotlib
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting gymnasium[box2d]
  Downloading gymnasium-0.28.1-py3-none-any.whl (925 kB)
    925.5/925.5 KB 16.1 MB/s eta 0:00:00
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[box2d]) (1.22.4)
Collecting jax-jumpy>=1.0.0
  Downloading jax_jumpy-1.0.0-py3-none-any.whl (20 kB)
Requirement already satisfied: cloudpickle>=1.2.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[box2d]) (2.2.1)
Requirement already satisfied: importlib-metadata>=4.8.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[box2d]) (6.1.0)
Requirement already satisfied: typing-extensions>=4.3.0 in /usr/local/lib/python3.9/dist-packages (from gymnasium[box2d]) (4.5.0)
Collecting farama-notifications>=0.0.1
  Downloading Farama_Notifications-0.0.4-py3-none-any.whl (2.5 kB)
Collecting pygame==2.1.3
  Downloading pygame-2.1.3-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (13.7 MB)
    13.7/13.7 MB 58.8 MB/s eta 0:00:00
Collecting box2d-py==2.3.5
  Downloading box2d-py-2.3.5.tar.gz (374 kB)
    374.4/374.4 KB 14.0 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
Collecting swig==4.*
  Downloading swig-4.1.1-py2.py3-none-manylinux_2_5_x86_64.manylinux1_x86_64.whl (1.8 MB)
    1.8/1.8 MB 47.9 MB/s eta 0:00:00
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.9/dist-packages (from importlib-metadata>=4.8.0->gymnasium[box2d])
Building wheels for collected packages: box2d-py
  error: subprocess-exited-with-error

  × python setup.py bdist_wheel did not run successfully.
  | exit code: 1
  | See above for output.

  note: This error originates from a subprocess, and is likely not a problem with pip.
Building wheel for box2d-py (setup.py) ... error
ERROR: Failed building wheel for box2d-py
Running setup.py clean for box2d-py
Failed to build box2d-py
Installing collected packages: swig, farama-notifications, box2d-py, pygame, jax-jumpy, gymnasium
Running setup.py install for box2d-py ... done
DEPRECATION: box2d-py was installed using the legacy 'setup.py install' method, because a wheel could not be built for it. A possib
Attempting uninstall: pygame
  Found existing installation: pygame 2.3.0
  Uninstalling pygame-2.3.0:
    Successfully uninstalled pygame-2.3.0
Successfully installed box2d-py-2.3.5 farama-notifications-0.0.4 gymnasium-0.28.1 jax-jumpy-1.0.0 pygame-2.1.3 swig-4.1.1
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: matplotlib in /usr/local/lib/python3.9/dist-packages (3.7.1)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.9/dist-packages (from matplotlib) (3.0.9)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.9/dist-packages (from matplotlib) (23.0)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.9/dist-packages (from matplotlib) (4.39.3)
Requirement already satisfied: importlib-resources>=3.2.0 in /usr/local/lib/python3.9/dist-packages (from matplotlib) (5.12.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.9/dist-packages (from matplotlib) (1.0.7)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.9/dist-packages (from matplotlib) (0.11.0)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.9/dist-packages (from matplotlib) (8.4.0)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.9/dist-packages (from matplotlib) (2.8.2)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.9/dist-packages (from matplotlib) (1.4.4)
Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.9/dist-packages (from matplotlib) (1.22.4)
Requirement already satisfied: zipp>=3.1.0 in /usr/local/lib/python3.9/dist-packages (from importlib-resources>=3.2.0->matplotlib) (3
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.9/dist-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)
```

```
1 import gymnasium as gym
2 import random
```

```

3 import torch
4 import torch.nn as nn
5 import torch.nn.functional as F
6 import torch.optim as optim
7 import argparse
8 import numpy as np
9 from collections import deque, namedtuple
10 import matplotlib.pyplot as plt
11 import base64, io
12
13 # For visualization
14 from gymnasium.wrappers.monitoring import video_recorder
15 from IPython.display import HTML
16 from IPython import display
17 import glob

1 class ExperienceReplay:
2     """
3     Based on the Replay Buffer implementation of TD3
4     Reference: https://github.com/sfujim/TD3/blob/master/utils.py
5     """
6     def __init__(self, state_dim, action_dim, max_size, batch_size, gpu_index=0):
7         self.max_size = max_size
8         self.ptr = 0
9         self.size = 0
10        self.state = np.zeros((max_size, state_dim))
11        self.action = np.zeros((max_size, action_dim))
12        self.next_state = np.zeros((max_size, state_dim))
13        self.reward = np.zeros((max_size, 1))
14        self.done = np.zeros((max_size, 1))
15        self.batch_size = batch_size
16        self.device = torch.device('cuda', index=gpu_index) if torch.cuda.is_available() else torch.device('cpu')
17
18
19    def add(self, state, action, reward, next_state, done):
20        self.state[self.ptr] = state
21        self.action[self.ptr] = action
22        self.next_state[self.ptr] = next_state
23        self.reward[self.ptr] = reward
24        self.done[self.ptr] = done
25        self.ptr = (self.ptr + 1) % self.max_size
26        self.size = min(self.size + 1, self.max_size)
27
28    def sample(self):
29        ind = np.random.randint(0, self.size, size=self.batch_size)
30
31        return (
32            torch.FloatTensor(self.state[ind]).to(self.device),
33            torch.FloatTensor(self.action[ind]).long().to(self.device),
34            torch.FloatTensor(self.reward[ind]).to(self.device),
35            torch.FloatTensor(self.next_state[ind]).to(self.device),
36            torch.FloatTensor(self.done[ind]).to(self.device)
37        )
38
39
40
41 class QNetwork(nn.Module):
42     """
43     Q Network: designed to take state as input and give out Q values of actions as output
44     """
45
46    def __init__(self, state_dim, action_dim):
47        """
48        state_dim (int): state dimension
49        action_dim (int): action dimension
50        """
51        super(QNetwork, self).__init__()
52        self.l1 = nn.Linear(state_dim, 64)
53        self.l2 = nn.Linear(64, 64)
54        self.l3 = nn.Linear(64, action_dim)
55
56    def forward(self, state):
57        q = F.relu(self.l1(state))
58        q = F.relu(self.l2(q))
59        return self.l3(q)

```

```

1 class DQNAgent():
2
3     def __init__(self,
4         state_dim,
5         action_dim,
6         discount=0.99,
7         tau=1e-3,
8         lr=5e-4,
9         update_freq=4,
10        max_size=int(1e5),
11        batch_size=64,
12        gpu_index=0
13    ):
14        """
15        state_size (int): dimension of each state
16        action_size (int): dimension of each action
17        discount (float): discount factor
18        tau (float): used to update q-target
19        lr (float): learning rate
20        update_freq (int): update frequency of target network
21        max_size (int): experience replay buffer size
22        batch_size (int): training batch size
23        gpu_index (int): GPU used for training
24        """
25        self.state_dim = state_dim
26        self.action_dim = action_dim
27        self.discount = discount
28        self.tau = tau
29        self.lr = lr
30        self.update_freq = update_freq
31        self.max_size = max_size
32        self.batch_size = batch_size
33        self.device = torch.device('cuda', index=gpu_index) if torch.cuda.is_available() else torch.device('cpu')
34
35
36        # Setting up the NNs
37        self.Q = QNetwork(state_dim, action_dim).to(self.device)
38        self.Q_target = QNetwork(state_dim, action_dim).to(self.device)
39        self.optimizer = optim.Adam(self.Q.parameters(), lr=self.lr)
40
41        # Experience Replay Buffer
42        self.memory = ExperienceReplay(state_dim,1,max_size,self.batch_size,gpu_index)
43
44        self.t_train = 0
45
46    def step(self, Exp_rep, targ_net, state, action, reward, next_state, done):
47        """
48        1. Adds (s,a,r,s') to the experience replay buffer, and updates the networks
49        2. Learns when the experience replay buffer has enough samples
50        3. Updates target network
51        """
52        self.t_train += 1
53        self.memory.add(state, action, reward, next_state, done)
54
55
56        # Experience Reply and Target Network Enabled
57        if Exp_rep and targ_net:
58
59            # Conducting Experience Replay
60            if self.memory.size > self.batch_size:
61                experiences = self.memory.sample()
62                self.learn(experiences, self.discount) #To be implemented
63
64            # Updating Target Network
65            if(self.t_train % self.update_freq) == 0:
66                self.target_update(self.Q, self.Q_target, self.tau)
67
68        # ONLY conducting Experience Replay
69        elif Exp_rep and not targ_net:
70
71            if self.memory.size > self.batch_size:
72                self.memory.add(state, action, reward, next_state, done)
73                experiences = self.memory.sample()
74                self.learn_exp_rep(experiences, self.discount)
75            # Target Network should not be updated
76
77        # ONLY updating Target Network

```

```

78         else:
79
80             if self.memory.size > self.batch_size:
81                 self.learn_targ_net(state, action, reward, next_state, done, self.discount)
82                 if(self.t_train % self.update_freq) == 0:
83                     self.target_update(self.Q, self.Q_target, self.tau)
84
85 #To be implemented
86
87 def select_action(self, state, epsilon = 0.):
88     if np.random.random() > epsilon:
89
90         state = torch.tensor(state, dtype=torch.float32, device=self.device) # converting our state to pytorch tensor
91         #self.Q(state) # writing updated tensor to device
92         self.Q.eval()
93         with torch.no_grad():
94             actions=self.Q.forward(state)
95         self.Q.train()
96         action = torch.argmax(actions).item() # item() changes from tensor to integer
97     else:
98         action=np.random.choice(self.action_dim)
99
100    return action
101
102 def learn(self, experiences, discount):
103     """
104     TODO: Complete this block to update the Q-Network using the target network
105     1. Compute target using self.Q_target ( tar get = r + discount * max_b [Q_target(s,b)] )
106     2. Compute Q(s,a) using self.Q
107     3. Compute MSE loss between step 1 and step 2
108     4. Update your network
109     Input: experiences consisting of states,actions,rewards,next_states and discount factor
110     Return: None
111     """
112     states, actions, rewards, next_states, dones = experiences
113
114
115     q_eval = self.Q(states).gather(1, actions)
116     q_next = self.Q_target(next_states).detach().max(1)[0].unsqueeze(1)
117
118     # Calculating q_target
119     q_target = rewards + discount * q_next * (1-dones)
120
121     # Calculating loss and backpropogating
122     loss = F.mse_loss(q_eval, q_target)
123     self.optimizer.zero_grad()
124     loss.backward()
125     self.optimizer.step()
126     self.target_update(self.Q, self.Q_target, self.tau)
127
128
129 def learn_exp_rep(self, experiences, discount):
130     """
131     TODO: Complete this block to update the Q-Network using the target network
132     1. Compute target using self.Q_target ( tar get = r + discount * max_b [Q_target(s,b)] )
133     2. Compute Q(s,a) using self.Q
134     3. Compute MSE loss between step 1 and step 2
135     4. Update your network
136     Input: experiences consisting of states,actions,rewards,next_states and discount factor
137     Return: None
138     """
139     states, actions, rewards, next_states, dones = experiences
140
141
142     q_eval = self.Q(states).gather(1, actions)
143     q_next = self.Q_target(next_states).detach().max(1)[0].unsqueeze(1)
144     q_target = rewards + discount * q_next * (1-dones)
145     loss = F.mse_loss(q_eval, q_target)
146     self.optimizer.zero_grad()
147     loss.backward()
148     self.optimizer.step()
149     #self.target_update(self.Q, self.Q_target, self.tau)
150
151 def learn_targ_net(self, states, actions, rewards, next_states, dones, discount):
152     """
153     TODO: Complete this block to update the Q-Network using the target network
154     1. Compute target using self.Q_target ( tar get = r + discount * max_b [Q_target(s,b)] )
155     2. Compute Q(s,a) using self.Q

```

```

155     3. Compute MSE loss between step 1 and step 2
156     4. Update your network
157     Input: experiences consisting of states,actions,rewards,next_states and discount factor
158     Return: None
159     """
160     # Storing states
161
162     states, actions, rewards, next_states, dones = [states], [actions], [rewards], [next_states], [dones]
163     states = torch.tensor(states).float().to(self.device)
164     actions = torch.tensor(actions).long().unsqueeze(1).to(self.device) # Fix here
165     rewards = torch.tensor(rewards).float().unsqueeze(1).to(self.device) # Fix here
166     next_states = torch.tensor(next_states).float().to(self.device)
167     dones = torch.tensor(dones).unsqueeze(1).float().to(self.device)
168
169     # Calculating q_target
170     q_eval = self.Q(states).gather(1, actions)
171     q_next = self.Q_target(next_states).detach().max(1)[0].unsqueeze(1)
172
173     # Calculating loss and backpropogating
174     q_target = rewards + discount * q_next * (1-dones)
175     loss = F.mse_loss(q_eval, q_target)
176     self.optimizer.zero_grad()
177     loss.backward()
178     self.optimizer.step()
179     self.target_update(self.Q, self.Q_target, self.tau)
180
181     def target_update(self, Q, Q_target, tau):
182         """
183         TODO: Update the target network parameters (param_target) using current Q parameters (param_Q)
184         Perform the update using tau, this ensures that we do not change the target network drastically
185         1. param_target = tau * param_Q + (1 - tau) * param_target
186         Input: Q,Q_target,tau
187         Return: None
188         """
189         ##### TYPE YOUR CODE HERE #####
190         for param_target, param_local in zip(Q_target.parameters(), Q.parameters()):
191             param_target.data.copy_(tau*param_local.data + (1.0-tau)*param_target.data)
192
193
194         #param_target = tau * param_Q + (1 - tau) * param_target

```

Now I'm comparing only target network and Experience Replay. \

This is the code for ONLY Experience Replay

```

1 def exp_rep_only():
2     seed = 0
3     n_episodes = 1500
4     batch_size = 64
5     discount = 0.99
6     lr = 5e-4 # learning rate
7     tau = 0.001 # soft update of target network
8     max_size = int(1e5)
9     update_freq = 4
10    gpu_index = 0
11    max_eps_len = 1000
12    #exploration strategy
13
14    epsilon = 1
15
16    # making the environment
17    env = gym.make("LunarLander-v2")
18
19    #setting seeds
20    torch.manual_seed(seed)
21    np.random.seed(seed)
22    random.seed(seed)
23
24    state_dim = env.observation_space.shape[0]
25    action_dim = env.action_space.n
26
27    kwargs = {
28        "state_dim":state_dim,
29        "action_dim":action_dim,
30        "discount":discount,
31        "tau":tau,


```

```

32     "lr":lr,
33     "update_freq":update_freq,
34
35     "max_size":max_size,
36     "batch_size":batch_size,
37     "gpu_index":gpu_index
38 }
39 exp_rep_only_learner = DQNAgent(**kwargs) #Creating the DQN learning agent
40 moving_window = deque(maxlen=100)
41 reward_store = []
42 print("Experience Replay Only")
43
44
45 for e in range(n_episodes):
46     state, _ = env.reset(seed=seed)
47     curr_reward = 0
48     for t in range(max_eps_len):
49         action = exp_rep_only_learner.select_action(state, epsilon) #To be implemented
50         n_state,reward,terminated,truncated,_ = env.step(action)
51         done = terminated or truncated
52         exp_rep_only_learner.step(True, False, state, action, reward, n_state, done) #To be implemented
53         state = n_state
54         curr_reward += reward
55         if done:
56             break
57     moving_window.append(curr_reward)
58     reward_store.append(curr_reward)
59
60
61     """
62     TODO: Write code for decaying the exploration rate using args.epsilon_decay
63     and args.epsilon_end. Note that epsilon has been initialized to args.epsilon_start
64     1. You are encouraged to try new methods
65     """
66     ##### TYPE YOUR CODE HERE #####
67     #####
68     plt.plot(np.arange(len(reward_store)), reward_store)
69     plt.ylabel('Score')
70     plt.xlabel('Episode')
71     plt.title('Decaying Epsilon DQN')
72     plt.show()
73
74     if e % 100 == 0:
75         print('Episode Number {} Average Episodic Reward (over 100 episodes): {:.2f}'.format(e, np.mean(moving_window)))
76
77
78     ##### TYPE YOUR CODE HERE #####
79     #####
80     averages = []
81     window_size=50
82
83     for i in range(len(reward_store) - window_size + 1):
84         window = reward_store[i:i+window_size]
85         average = sum(window) / window_size
86         averages.append(average)
87
88     plt.plot(averages, color='g')
89     plt.ylabel('Score')
90     plt.xlabel('Episodes')
91     plt.title('Only Experience Replay DQN')
92     plt.show()
93
94 exp_rep_only()

```

Experience Replay Only



This is the code for only Target Network

```

1  def only_tar_net_trainer():
2      seed = 0
3      n_episodes = 500
4      batch_size = 64
5      discount = 0.99
6      lr = 5e-4
7      tau = 0.001
8      max_size = int(1e5)
9      update_freq = 4
10     gpu_index = 0
11     max_eps_len = 1000
12
13     #exploration strategy
14
15     # making the environment
16     env = gym.make("LunarLander-v2")
17     #setting seeds
18     torch.manual_seed(seed)
19     np.random.seed(seed)
20     random.seed(seed)
21     state_dim = env.observation_space.shape[0]
22     action_dim = env.action_space.n
23     kwargs = {
24         "state_dim":state_dim,
25         "action_dim":action_dim,
26         "discount":discount,
27         "tau":tau,
28         "lr":lr,
29         "update_freq":update_freq,
30         "max_size":max_size,
31         "batch_size":batch_size,
32         "gpu_index":gpu_index
33     }
34     disabled_decaying_learner = DQNAgent(**kwargs) #Creating the DQN learning agent
35     moving_window = deque(maxlen=100)
36
37     reward_store = []
38

```



```

39     print("Target Network DQN")
40
41     epsilon_start = 1          # start value of epsilon
42     epsilon_end = 0.01        # end value of epsilon
43     epsilon_decay = 0.995     # decay value of epsilon
44
45     epsilon_by_step = lambda step: float(epsilon_end+(epsilon_start - epsilon_end)*np.exp(-1. * step / epsilon_decay))
46
47
48     for e in range(n_episodes):
49         state, _ = env.reset(seed=seed)
50         curr_reward = 0
51         for t in range(max_eps_len):
52             action = disabled_decaying_learner.select_action(state, epsilon_by_step(t)) #To be implemented
53             n_state, reward, terminated, truncated, _ = env.step(action)
54             done = terminated or truncated
55             disabled_decaying_learner.step(False, True, state, action, reward, n_state, done) #To be implemented
56             state = n_state
57             curr_reward += reward
58             if done:
59                 break
60         reward_store.append(curr_reward)
61         moving_window.append(curr_reward)
62
63         """
64         TODO: Write code for decaying the exploration rate using args.epsilon_decay
65         and args.epsilon_end. Note that epsilon has been initialized to args.epsilon_start
66         1. You are encouraged to try new methods
67         """
68         ##### TYPE YOUR CODE HERE #####
69         #####
70
71         if e % 100 == 0:
72             print('Episode Number {} Average Episodic Reward (over 100 episodes): {:.2f}'.format(e, np.mean(moving_window)))
73
74         """
75         TODO: Write code for
76         1. Logging and plotting
77         2. Rendering the trained agent
78         """
79         ##### TYPE YOUR CODE HERE #####
80         #####
81         plt.plot(np.arange(len(reward_store)), reward_store)
82         plt.ylabel('Score')
83         plt.xlabel('Episode')
84         plt.title('Decaying Epsilon DQN')
85         plt.show()
86
87         plt.plot(np.arange(len(moving_window)), moving_window)
88         plt.ylabel('Average')
89         plt.xlabel('Episodes')
90         plt.title('Only Target Network DQN')
91         plt.show()
92
93     only_tar_net_trainer()

```


Target Network DQN

Episode Number 0 Average Episodic Reward (over 100 episodes): -121.30
Episode Number 100 Average Episodic Reward (over 100 episodes): -246.68
Episode Number 200 Average Episodic Reward (over 100 episodes): -170.87
Episode Number 300 Average Episodic Reward (over 100 episodes): -64.01
Episode Number 400 Average Episodic Reward (over 100 episodes): -152.57



[Colab paid products](#) - [Cancel contracts here](#)

✓ 6m 12s completed at 11:03 PM

● ×