

A5_Q4_Deshmane_Aakash

April 27, 2023

AAKASH DESHMANE UIN : 133008022 ECEN 743 ASSIGNMENT 5 POLICY GRADIENT IMPLEMENTATION OF MOUNTAIN CAR CONTINUOUS

```
[ ]: import gymnasium as gym
import random
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import argparse
import numpy as np
import math
from collections import deque
import matplotlib.pyplot as plt
```

VALUE NETWORK

```
[ ]: class value_network(nn.Module):
    """
    Value Network: Designed to take in state as input and give value as
    ↪output
    Used as a baseline in Policy Gradient (PG) algorithms
    """
    def __init__(self, state_dim):
        """
        state_dim (int): state dimension
        """
        super(value_network, self).__init__()
        self.l1 = nn.Linear(state_dim, 128)
        self.l2 = nn.Linear(128, 64)
        self.l3 = nn.Linear(64, 1)

    def forward(self, state):
        """
        Input: State
        Output: Value of state
        """
        v = F.tanh(self.l1(state))
```

```

v = F.tanh(self.l2(v))
return self.l3(v)

```

POLICY NETWORK

```

[ ]: class policy_network(nn.Module):
    """
    Policy Network: Designed for continuous action space, where given a
    state, the network outputs the mean and standard deviation of the action
    """
    def __init__(self, state_dim, action_dim, log_std = 0.0):
        """
        state_dim (int): state dimension
        action_dim (int): action dimension
        log_std (float): log of standard deviation (std)
        """
        super(policy_network, self).__init__()
        self.state_dim = state_dim
        self.action_dim = action_dim
        self.l1 = nn.Linear(state_dim, 64)
        self.l2 = nn.Linear(64, 64)
        #self.l3 = nn.Linear(64, 64)
        #self.l4 = nn.Linear(64, 64)
        self.mean = nn.Linear(64, action_dim)
        self.log_std = nn.Parameter(torch.ones(1, action_dim) * log_std)

    def forward(self, state):
        """
        Input: State
        Output: Mean, log_std and std of action
        """
        a = F.tanh(self.l1(state))
        a = F.tanh(self.l2(a))
        a_mean = self.mean(a)
        a_log_std = self.log_std.expand_as(a_mean)
        a_std = torch.exp(a_log_std)
        return a_mean, a_log_std, a_std

    def select_action(self, state):
        """
        Input: State
        Output: Sample drawn from a normal distribution with mean and std
        """

        a_mean, _, a_std = self.forward(state)

```

```

        action = torch.normal(a_mean)
        return action

    def get_log_prob(self, state, action):
        '''
        Input: State, Action
        Output: log probabilities
        '''
        mean, log_std, std = self.forward(state)
        var = std.pow(2)
        log_density = -(action - mean).pow(2) / (2 * var) - 0.5 * math.
        ↪ log(2 * math.pi) - log_std
        return log_density.sum(1, keepdim=True)

```

```

[ ]: class PGAgent():
    '''
    An agent that performs different variants of the PG algorithm
    '''
    def __init__(self,
        state_dim,
        action_dim,
        discount=0.99,
        lr=1e-3,
        gpu_index=0,
        seed=0,
        env="LunarLander-v2"
    ):
        """
        state_size (int): dimension of each state
        action_size (int): dimension of each action
        discount (float): discount factor
        lr (float): learning rate
        gpu_index (int): GPU used for training
        seed (int): Seed of simulation
        env (str): Name of environment
        """
        self.state_dim = state_dim
        self.action_dim = action_dim
        self.discount = discount
        self.lr = lr
        self.device = torch.device('cuda', index=gpu_index) if torch.
        ↪ cuda.is_available() else torch.device('cpu')
        self.env_name = env
        self.seed = seed
        self.policy = policy_network(state_dim, action_dim)
        self.value = value_network(state_dim)

```

```

        self.optimizer_policy = torch.optim.Adam(self.policy.
↳parameters(), lr=self.lr)
        self.optimizer_value = torch.optim.Adam(self.value.
↳parameters(), lr=self.lr)

    def sample_traj(self, batch_size=4000, evaluate = False):
        '''
        Input:
            batch_size: minimum batch size needed for update
            evaluate: flag to be set during evaluation
        Output:
            states, actions, rewards, not_dones, episodic reward_
↳
        '''
        self.policy.to("cpu") #Move network to CPU for sampling
        env = gym.make(self.env_name)
        states = []
        actions = []
        rewards = []
        n_dones = []
        curr_reward_list = []
        while len(states) < batch_size:
            state, _ = env.reset(seed=self.seed)
            curr_reward = 0
            for t in range(1000):
                state_ten = torch.from_numpy(state).float().
↳unsqueeze(0)

                with torch.no_grad():
                    if evaluate:
                        action = self.
↳policy(state_ten)[0][0].numpy() # Take mean action during evaluation
                    else:
                        action = self.policy.
↳select_action(state_ten)[0].numpy() # Sample from distribution during_
↳training

                action = action.astype(np.float32)
                n_state, reward, terminated, truncated, _ = env.
↳step(action) # Execute action in the environment

                done = terminated or truncated
                states.append(state)
                actions.append(action)
                rewards.append(reward)
                n_done = 0 if done else 1
                n_dones.append(n_done)
                state = n_state
                curr_reward += reward

```

```

        if done:
            break
        curr_reward_list.append(curr_reward)
    if evaluate:
        return np.mean(curr_reward_list)
    return states, actions, rewards, n_dones, np.mean(curr_reward_list)

def update(self, states, actions, rewards, n_dones, update_type='Baseline'):

    self.policy.to(self.device) #Move policy to GPU
    if update_type == "Baseline":
        self.value.to(self.device) #Move value to GPU
        states_ten = torch.from_numpy(np.stack(states)).to(self.device)
        ↪ #Convert to tensor and move to GPU
        action_ten = torch.from_numpy(np.stack(actions)).to(self.
        ↪device) #Convert to tensor and move to GPU
        rewards_ten = torch.from_numpy(np.stack(rewards)).to(self.
        ↪device) #Convert to tensor and move to GPU
        n_dones_ten = torch.from_numpy(np.stack(n_dones)).to(self.
        ↪device) #Convert to tensor and move to GPU

    if update_type == "Rt":

        rt = torch.zeros(rewards_ten.shape[0],1).to(self.device)
        rt_accum = 0
        for t in reversed(range(rewards_ten.shape[0])):
            rt_accum = rewards_ten[t] + rt_accum *
            ↪self.discount * n_dones_ten[t]
            rt[t] = rt_accum

        log_prob = self.policy.get_log_prob(states_ten,
        ↪action_ten)

        policy_loss = -(log_prob * rt.detach()).mean()

        self.optimizer_policy.zero_grad()
        policy_loss.backward()
        self.optimizer_policy.step()

    if update_type == 'Gt':
        gt = torch.zeros(rewards_ten.shape[0],1).to(self.device)
        g = 0
        for i in reversed(range(rewards_ten.size(0))):

```

```

        g = rewards_ten[i] + self.discount * g
    ↪*(n_dones_ten[i])

        gt[i] = g

    gt = (gt - gt.mean()) / gt.std() #Helps with learning
    ↪stablility

    log_prob = self.policy.get_log_prob(states_ten,
    ↪action_ten)

    policy_loss = -(log_prob * gt.detach()).mean()

    self.optimizer_policy.zero_grad()
    policy_loss.backward()
    self.optimizer_policy.step()

    if update_type == 'Gt_with_Baseline':
        '''
        TODO: Peform PG using reward_to_go and baseline
        1. Compute values of states, this will be used as the
        ↪baseline

        2. Compute reward_to_go (gt) using rewards_ten and
        ↪n_dones_ten

        3. gt should be of the same length as rewards_ten
        4. Compute advantages
        5. Update the value network to predict gt for each
        ↪state (L2 norm)

        6. Compute log probabilities using states_ten and
        ↪action_ten

        7. Compute policy loss (using advantages) and update
        ↪the policy

        '''
        state_t = torch.FloatTensor(states).to(self.device)

        # STEP 1 CALCULATE VALUES
        with torch.no_grad():
            self.value.to(self.device)
            val = self.value(states_ten).to(self.
            ↪device)

        # gt SHOULD HAVE THE SAME LENGTH AS rewards_ten
        gt = torch.zeros(rewards_ten.shape[0],1).to(self.device)

        g=0

        # STEP 2 : COMPUTE REWARD-TO-GO (gt) and ADVANTAGES
        returns = torch.zeros((rewards_ten.shape[0], 1)).
        ↪to(self.device)

```

```

advantages = torch.zeros((rewards_ten.shape[0], 1)).
↳to(self.device)
s = rewards_ten.size(0)
for i in reversed(range(s)):
    g = rewards_ten[i] + self.discount * g *
↳n_dones_ten[i]
    gt[i] = g

# STEP 4 : COMPUTE ADVANTAGES
advantages = gt - val

# Normalize advantages
advantages = (advantages - advantages.mean()) /

↳advantages.std()

# STEP 5 : UPDATE VALUE NETWORK TO PREDICT gt FOR EACH
↳STATE (L2 NORM)
loss = torch.nn.MSELoss()
value_loss = loss(self.value(states_ten), gt)
self.optimizer_value.zero_grad()
value_loss.backward()
self.optimizer_value.step()

# STEP 6 : COMPUTE LOG PROBABILITIES USING states_ten
↳and Compute log probabilities using states_ten and action_ten
log_probs = self.policy.get_log_prob(states_ten,
↳action_ten)

# STEP 7 : COMPUTE POLICY LOSS AND UPDATE POLICY
self.optimizer_policy.zero_grad()
l = log_probs * advantages.detach()
loss = -(l).mean()
loss.backward()
self.optimizer_policy.step()

#
↳

```

RUNNING MOUNTAIN CAR FOR MORE NUMBER OF ITERATIONS TO ACHIEVE CONVERGENCE
HYPERPARAMETERS TUNED TO: * ITERATIONS : 800 * DISCOUNT : 0.99 *
BATCH SIZE : 4000 * LEARNING RATE : 8e-4

```

[ ]: env_type = "MountainCarContinuous-v0" # Gymnasium environment name
seed=0 # Sets Gym, PyTorch and Numpy seeds
n_iter = 800 # Maximum number of training iterations
discount = 0.99 # Discount factor
batch_size = 4000 # Training samples in each batch of training

```

```

lr = 8e-3          # Learning rate
gpu_index = 0      # GPU index
algo = "Gt_with_Baseline"          # PG algorithm type.
    ↳ Baseline_with_Gt/Gt/Rt

# Making the environment
env = gym.make(env_type)

# Setting seeds
torch.manual_seed(seed)
np.random.seed(seed)
random.seed(seed)

state_dim = env.observation_space.shape[0]
print(state_dim)
action_dim = env.action_space.shape[0]

kwargs = {
    "state_dim":state_dim,
    "action_dim":action_dim,
    "discount":discount,
    "lr":lr,
    "gpu_index":gpu_index,
    "seed":seed,
    "env":env_type
}
learner = PGAgent(**kwargs) # Creating the PG learning agent
average_rewards=[]
moving_window = deque(maxlen=10)
old_reward=-1
for e in range(n_iter):
    states,actions,rewards,n_dones,train_reward = learner.
    ↳ sample_traj(batch_size=batch_size)
    learner.update(states,actions,rewards,n_dones,algo)
    eval_reward= learner.sample_traj(evaluate=True)
    moving_window.append(eval_reward)
    if not e%100: print('Training Iteration {} Training Reward: {:.2f}
    ↳ Evaluation Reward: {:.2f} \
        Average Evaluation Reward: {:.2f}'.format(e,train_reward,eval_reward,np.
    ↳ mean(moving_window)))

    average_rewards.append(np.mean(moving_window))

    if np.mean(moving_window) > old_reward:
        old_reward = np.mean(moving_window)

```



```

        torch.save(learner.policy.state_dict(), (algo + '_checkpoint1.
        ↪pth'))

window_size = 20
averages = []
fig1=plt.figure()
for i in range(len(average_rewards)-window_size + 1):
    window = average_rewards[i:i+window_size]
    average = sum(window)/window_size
    averages.append(average)
plt.plot(averages)
plt.show()
fig2=plt.figure()

plt.plot(average_rewards)
plt.ylabel('Episodic Cumulative Reward')
plt.xlabel('Episode #')
plt.title('Curve for Episodic Cumulative Reward for algorithm = {}'.
        ↪format(algo))
plt.show()

```

2

/tmp/ipykernel_16353/4280312714.py:122: UserWarning: Creating a tensor from a list of numpy.ndarrays is extremely slow. Please consider converting the list to a single numpy.ndarray with numpy.array() before converting to a tensor.

(Triggered internally at ../torch/csrc/utils/tensor_new.cpp:245.)

```
state_t = torch.FloatTensor(states).to(self.device)
```

Training Iteration 0 Training Reward: -102.77 Evaluation Reward: -63.52

Average Evaluation Reward: -63.52

Training Iteration 100 Training Reward: -94.95 Evaluation Reward: -0.19

Average Evaluation Reward: -0.61

Training Iteration 200 Training Reward: -101.77 Evaluation Reward: -1.52

Average Evaluation Reward: -0.89

Training Iteration 300 Training Reward: 53.55 Evaluation Reward: 75.41 Average

Evaluation Reward: 57.69

Training Iteration 400 Training Reward: 55.07 Evaluation Reward: 67.36 Average

Evaluation Reward: 81.26

Training Iteration 500 Training Reward: 75.35 Evaluation Reward: 82.92 Average

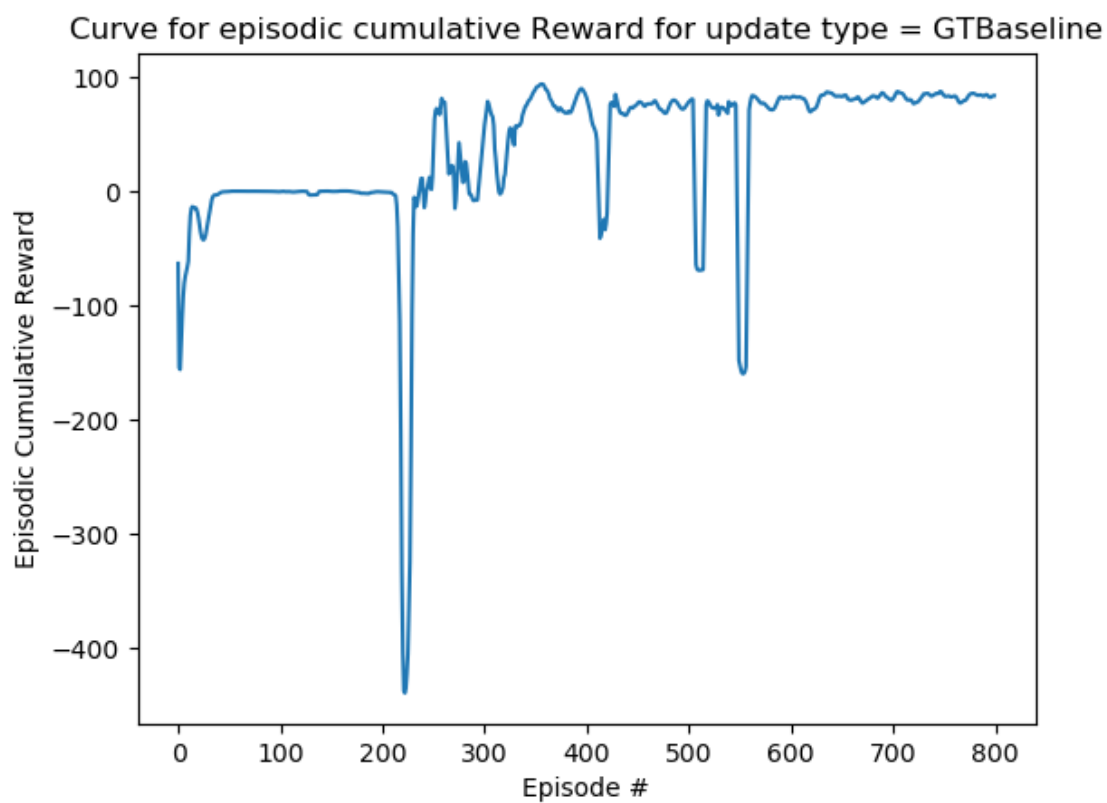

Evaluation Reward: 78.08

Training Iteration 600 Training Reward: 70.57 Evaluation Reward: 78.63 Average

Evaluation Reward: 81.68

Training Iteration 700 Training Reward: 72.33 Evaluation Reward: 91.86 Average

Evaluation Reward: 80.78



```
[ ]:
```

THE CODE BELOW IS JUST TO EXPORT THE VIDEO AND DOES NOT TAKE PART IN THE ALGORITHM

```
[ ]: #For visualization
import gymnasium as gym
from gym.wrappers.monitoring import video_recorder
from IPython.display import HTML
from IPython import display
import glob
import cv2
```

VIDEO FUNCTION

```
[ ]: def video_fn(agent, env_name, algo):
    env = gym.make(env_name, render_mode="rgb_array")
    fourcc = cv2.VideoWriter_fourcc(*'mp4v')
    video = cv2.VideoWriter(algo+'_video.mp4', fourcc, 30, (600, 400))
    agent.policy.load_state_dict(torch.load(algo+"_checkpoint.pth"))
    agent.policy.eval()
    state, _ = env.reset()
    done = False
    while not done:
        frame = env.render()
        video.write(frame)
        state_ten = torch.from_numpy(state).float().unsqueeze(0)
        action = agent.policy.select_action(state_ten)[0].detach().numpy()
        action = action.astype(np.float64)
        n_state, reward, terminated, truncated, _ = env.step(action)
        done = terminated or truncated
        state = n_state
    env.close()
    video.release()
```

EXPORTING VIDEO

```
[ ]: env_type = "MountainCarContinuous"
env = gym.make(env_type)
state_dim = env.observation_space.shape[0]
action_dim = env.action_space.shape[0]
plotter_agent = PGAgent(state_dim, action_dim)
#video_fn(plotter_agent, "MountainCarContinuous", "Rt")
#video_fn(plotter_agent, "MountainCarContinuous", "Gt")
video_fn(plotter_agent, "MountainCarContinuous", "Gt_with_Baseline")
```